

Machine Learning Methology for Diagnosing Chronic Kidney Disease

All features

```
In [50]: #import necessary library
import pandas as pd # for data analysis
import numpy as np # for numerical calculation
import matplotlib.pyplot as plt # for data visualization
import seaborn as sns #for interactive data visualization
# Ignore warnings
import warnings
warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', None) # or use a very large number like 1000
```

```
In [2]: #Load data set
df = pd.read_csv("kidney_disease.csv")
#see forward data
df.head(5)
```

```
Out[2]:
```

	id	age	bp	sg	al	su	rbc	pc	pcc	ba	bgr	bu	sc	sod	pcv
0	0	48.0	80.0	1.020	1.0	0.0	NaN	normal	notpresent	notpresent	121.0	36.0	1.2	NaN	NaN
1	1	7.0	50.0	1.020	4.0	0.0	NaN	normal	notpresent	notpresent	NaN	18.0	0.8	NaN	NaN
2	2	62.0	80.0	1.010	2.0	3.0	normal	normal	notpresent	notpresent	423.0	53.0	1.8	NaN	NaN
3	3	48.0	70.0	1.005	4.0	0.0	normal	abnormal	present	notpresent	117.0	56.0	3.8	111.0	2
4	4	51.0	80.0	1.010	2.0	0.0	normal	normal	notpresent	notpresent	106.0	26.0	1.4	NaN	NaN

```
In [3]: # see data shape(number of rows and columns)
df.shape
```

```
Out[3]: (400, 26)
```

```
In [4]: #see columns
df.columns
```

```
Out[4]: Index(['id', 'age', 'bp', 'sg', 'al', 'su', 'rbc', 'pc', 'pcc', 'ba', 'bgr',
              'bu', 'sc', 'sod', 'pot', 'hemo', 'pcv', 'wc', 'rc', 'htn', 'dm', 'cad',
              'appet', 'pe', 'ane', 'classification'],
              dtype='object')
```

```
In [5]: #dropping id column
df.drop('id', axis = 1, inplace = True)
```

```
In [6]: #rename columns name
df.columns = ['age', 'blood_pressure', 'specific_gravity', 'albumin', 'sugar', 'red_blood_cells', 'pus_cell_clumps', 'bacteria', 'blood_glucose_random', 'blood_urea', 'serum_creatinine', 'sodium', 'potassium', 'haemoglobin', 'packed_cell_volume', 'white_blood_cell_count', 'hypertension', 'diabetes_mellitus', 'coronary_artery_disease', 'appetite', 'peda_edema', 'aanemia', 'class']

df.columns
```

```
Out[6]: Index(['age', 'blood_pressure', 'specific_gravity', 'albumin', 'sugar', 'red_blood_cells', 'pus_cell', 'pus_cell_clumps', 'bacteria', 'blood_glucose_random', 'blood_urea', 'serum_creatinine', 'sodium', 'potassium', 'haemoglobin', 'packed_cell_volume', 'white_blood_cell_count', 'red_blood_cell_count', 'hypertension', 'diabetes_mellitus', 'coronary_artery_disease', 'appetite', 'peda_edema', 'aanemia', 'class'],
      dtype='object')
```

```
In [8]: # Basic statistical measurements
df.describe().T
```

```
Out[8]:
```

	count	mean	std	min	25%	50%	75%	max
age	391.0	51.483376	17.169714	2.000	42.00	55.00	64.50	90.000
blood_pressure	388.0	76.469072	13.683637	50.000	70.00	80.00	80.00	180.000
specific_gravity	353.0	1.017408	0.005717	1.005	1.01	1.02	1.02	1.025
albumin	354.0	1.016949	1.352679	0.000	0.00	0.00	2.00	5.000
sugar	351.0	0.450142	1.099191	0.000	0.00	0.00	0.00	5.000
blood_glucose_random	356.0	148.036517	79.281714	22.000	99.00	121.00	163.00	490.000
blood_urea	381.0	57.425722	50.503006	1.500	27.00	42.00	66.00	391.000
serum_creatinine	383.0	3.072454	5.741126	0.400	0.90	1.30	2.80	76.000
sodium	313.0	137.528754	10.408752	4.500	135.00	138.00	142.00	163.000
potassium	312.0	4.627244	3.193904	2.500	3.80	4.40	4.90	47.000
haemoglobin	348.0	12.526437	2.912587	3.100	10.30	12.65	15.00	17.800

```
In [11]: #see data info()
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                    391 non-null    float64
1   blood_pressure                        388 non-null    float64
2   specific_gravity                      353 non-null    float64
3   albumin                              354 non-null    float64
4   sugar                                 351 non-null    float64
5   red_blood_cells                       248 non-null    object
6   pus_cell                              335 non-null    object
7   pus_cell_clumps                       396 non-null    object
8   bacteria                              396 non-null    object
9   blood_glucose_random                  356 non-null    float64
10  blood_urea                            381 non-null    float64
11  serum_creatinine                      383 non-null    float64
12  sodium                                313 non-null    float64
13  potassium                             312 non-null    float64
14  haemoglobin                           348 non-null    float64
15  packed_cell_volume                    329 non-null    float64
16  white_blood_cell_count                294 non-null    float64
17  red_blood_cell_count                  269 non-null    float64
18  hypertension                          398 non-null    object
19  diabetes_mellitus                     398 non-null    object
20  coronary_artery_disease               398 non-null    object
21  appetite                              399 non-null    object
22  peda_edema                            399 non-null    object
23  aanemia                               399 non-null    object
24  class                                 400 non-null    object
dtypes: float64(14), object(11)
memory usage: 78.2+ KB
```

```
In [10]: # converting necessary columns to numerical type

df['packed_cell_volume'] = pd.to_numeric(df['packed_cell_volume'], errors='coerce')
df['white_blood_cell_count'] = pd.to_numeric(df['white_blood_cell_count'], errors='coerce')
df['red_blood_cell_count'] = pd.to_numeric(df['red_blood_cell_count'], errors='coerce')

In [12]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                   391 non-null    float64
1   blood_pressure                       388 non-null    float64
2   specific_gravity                     353 non-null    float64
3   albumin                             354 non-null    float64
4   sugar                                351 non-null    float64
5   red_blood_cells                      248 non-null    object
6   pus_cell                             335 non-null    object
7   pus_cell_clumps                      396 non-null    object
8   bacteria                             396 non-null    object
9   blood_glucose_random                 356 non-null    float64
10  blood_urea                           381 non-null    float64
11  serum_creatinine                     383 non-null    float64
12  sodium                               313 non-null    float64
13  potassium                            312 non-null    float64
14  haemoglobin                          348 non-null    float64
15  packed_cell_volume                   329 non-null    float64
16  white_blood_cell_count               294 non-null    float64
17  red_blood_cell_count                 269 non-null    float64
18  hypertension                         398 non-null    object
19  diabetes_mellitus                   398 non-null    object
20  coronary_artery_disease              398 non-null    object
21  appetite                             399 non-null    object
22  peda_edema                           399 non-null    object
23  aanemia                              399 non-null    object
24  class                                400 non-null    object
dtypes: float64(14), object(11)
memory usage: 78.2+ KB
```

Data Preprocessing

```
In [19]: # Extract categorical columns (object types)
categorical_columns = df.select_dtypes(include=['object']).columns.tolist()
print(categorical_columns)
print(len(categorical_columns))

['red_blood_cells', 'pus_cell', 'pus_cell_clumps', 'bacteria', 'hypertension', 'diabetes_mellitus', 'coronary_artery_disease', 'appetite', 'peda_edema', 'aanemia', 'class']
11
```

```
In [18]: # Extract numerical columns (integer and float types)
numerical_columns = df.select_dtypes(include=['int64', 'float64']).columns.tolist()
print(numerical_columns)
print(len(numerical_columns))

['age', 'blood_pressure', 'specific_gravity', 'albumin', 'sugar', 'blood_glucose_random', 'blood_urea', 'serum_creatinine', 'sodium', 'potassium', 'haemoglobin', 'packed_cell_volume', 'white_blood_cell_count', 'red_blood_cell_count']
14
```

```
In [20]: # Iterate through each categorical column
for column in categorical_columns:
    unique_values = df[column].unique()
    print(f"Unique values in '{column}': {unique_values}")
```

```

Unique values in 'red_blood_cells': [nan 'normal' 'abnormal']
Unique values in 'pus_cell': ['normal' 'abnormal' nan]
Unique values in 'pus_cell_clumps': ['notpresent' 'present' nan]
Unique values in 'bacteria': ['notpresent' 'present' nan]
Unique values in 'hypertension': ['yes' 'no' nan]
Unique values in 'diabetes_mellitus': ['yes' 'no' ' yes' '\tno' '\tyes' nan]
Unique values in 'coronary_artery_disease': ['no' 'yes' '\tno' nan]
Unique values in 'appetite': ['good' 'poor' nan]
Unique values in 'peda_edema': ['no' 'yes' nan]
Unique values in 'aanemia': ['no' 'yes' nan]
Unique values in 'class': ['ckd' 'ckd\t' 'notckd']

```

```

In [21]: # replace incorrect values

df['diabetes_mellitus'].replace(to_replace = {'\tno': 'no', '\tyes': 'yes', ' yes': 'yes'},

df['coronary_artery_disease'] = df['coronary_artery_disease'].replace(to_replace = '\t

df['class'] = df['class'].replace(to_replace = {'ckd\t': 'ckd', 'notckd': 'not ckd'})

```

```

In [22]: df['class'] = df['class'].map({'ckd': 0, 'not ckd': 1})
df['class'] = pd.to_numeric(df['class'], errors='coerce')
cols = ['diabetes_mellitus', 'coronary_artery_disease', 'class']

for col in cols:
    print(f"{col} has {df[col].unique()} values\n")

diabetes_mellitus has ['yes' 'no' nan] values

coronary_artery_disease has ['no' 'yes' nan] values

class has [0 1] values

```

```

In [23]: # Checking missing value
df.isnull().sum()

```

```
Out[23]:
```

age	9
blood_pressure	12
specific_gravity	47
albumin	46
sugar	49
red_blood_cells	152
pus_cell	65
pus_cell_clumps	4
bacteria	4
blood_glucose_random	44
blood_urea	19
serum_creatinine	17
sodium	87
potassium	88
haemoglobin	52
packed_cell_volume	71
white_blood_cell_count	106
red_blood_cell_count	131
hypertension	2
diabetes_mellitus	2
coronary_artery_disease	2
appetite	1
peda_edema	1
aanemia	1
class	0

dtype: int64

```
In [25]: # Checking missing values in numerical columns
numerical_columns = df.select_dtypes(include=['int64', 'float64']).columns.tolist()

for column in numerical_columns:
    missing_count = df[column].isnull().sum()
    print(f"Missing values in '{column}': {missing_count}")

Missing values in 'age': 9
Missing values in 'blood_pressure': 12
Missing values in 'specific_gravity': 47
Missing values in 'albumin': 46
Missing values in 'sugar': 49
Missing values in 'blood_glucose_random': 44
Missing values in 'blood_urea': 19
Missing values in 'serum_creatinine': 17
Missing values in 'sodium': 87
Missing values in 'potassium': 88
Missing values in 'haemoglobin': 52
Missing values in 'packed_cell_volume': 71
Missing values in 'white_blood_cell_count': 106
Missing values in 'red_blood_cell_count': 131
Missing values in 'class': 0
```

```
In [26]: # Checking missing values in categorical column
categorical_columns = df.select_dtypes(include=['object']).columns.tolist()

for column in categorical_columns:
    missing_count = df[column].isnull().sum()
    print(f"Missing values in '{column}': {missing_count}")
```

Missing values in 'red_blood_cells': 152
 Missing values in 'pus_cell': 65
 Missing values in 'pus_cell_clumps': 4
 Missing values in 'bacteria': 4
 Missing values in 'hypertension': 2
 Missing values in 'diabetes_mellitus': 2
 Missing values in 'coronary_artery_disease': 2
 Missing values in 'appetite': 1
 Missing values in 'peda_edema': 1
 Missing values in 'aanemia': 1

Handling Missing Value with KNNImputer

```
In [31]: from sklearn.impute import KNNImputer

# List of categorical and numerical columns
cat_cols = ['red_blood_cells', 'pus_cell', 'pus_cell_clumps', 'bacteria', 'hypertension', 'diabetes_mellitus', 'coronary_artery_disease', 'appetite', 'peda_edema', 'aanemia']
num_cols = ['age', 'blood_pressure', 'specific_gravity', 'albumin', 'sugar', 'blood_glucose_random', 'blood_urea', 'serum_creatinine', 'sodium', 'potassium', 'haemoglobin', 'packed_cell_volume', 'white_blood_cell_count', 'red_blood_cell_count']

# Initialize KNNImputer
knn_imputer = KNNImputer(n_neighbors=5) # You can adjust the number of neighbors

# Filling missing values for numerical columns
df[num_cols] = knn_imputer.fit_transform(df[num_cols])

# Filling missing values for categorical columns
df[cat_cols] = df[cat_cols].apply(lambda col: col.fillna(col.value_counts().idxmax()))

# Verify that missing values are filled
print(df.isnull().sum())
```

```
age                                0
blood_pressure                     0
specific_gravity                   0
albumin                           0
sugar                             0
red_blood_cells                    0
pus_cell                           0
pus_cell_clumps                    0
bacteria                           0
blood_glucose_random               0
blood_urea                         0
serum_creatinine                   0
sodium                             0
potassium                          0
haemoglobin                        0
packed_cell_volume                 0
white_blood_cell_count             0
red_blood_cell_count               0
hypertension                       0
diabetes_mellitus                  0
coronary_artery_disease            0
appetite                           0
peda_edema                         0
aanemia                           0
class                             0
dtype: int64
```

In []:

Exploratory data analysis

Numerical columns analysis

```
In [35]: import matplotlib.pyplot as plt

# Calculate the number of rows and columns needed for the grid
num_rows = len(num_cols) // 2 + len(num_cols) % 2
num_cols_per_row = 2

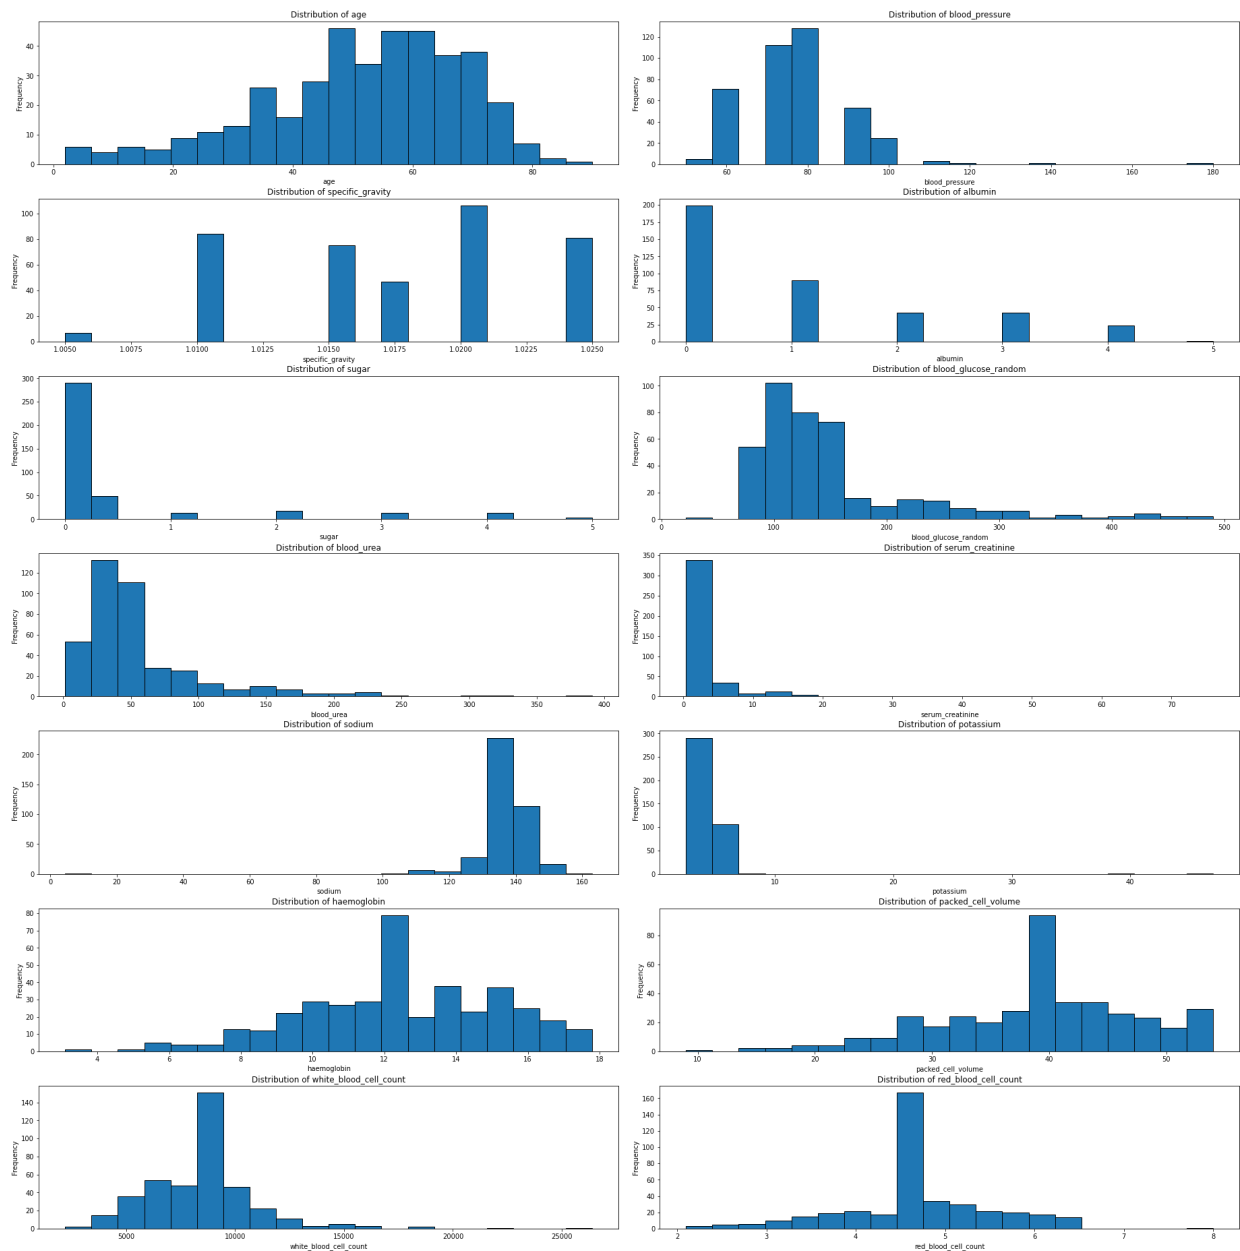
fig, axes = plt.subplots(num_rows, num_cols_per_row, figsize=(26, 26))
fig.tight_layout(pad=3.0) # Adjust the spacing between plots

for i, column in enumerate(num_cols):
    row = i // num_cols_per_row
    col = i % num_cols_per_row

    axes[row, col].hist(df[column], bins=20, edgecolor='k')
    axes[row, col].set_title(f'Distribution of {column}')
    axes[row, col].set_xlabel(column)
    axes[row, col].set_ylabel('Frequency')

# Hide any empty subplots
for i in range(len(num_cols), num_rows * num_cols_per_row):
    row = i // num_cols_per_row
    col = i % num_cols_per_row
    axes[row, col].axis('off')

plt.show()
```

```
In [39]: # Calculate the number of rows and columns needed for the grid
num_rows = len(num_cols) // 2 + len(num_cols) % 2
num_cols_per_row = 2

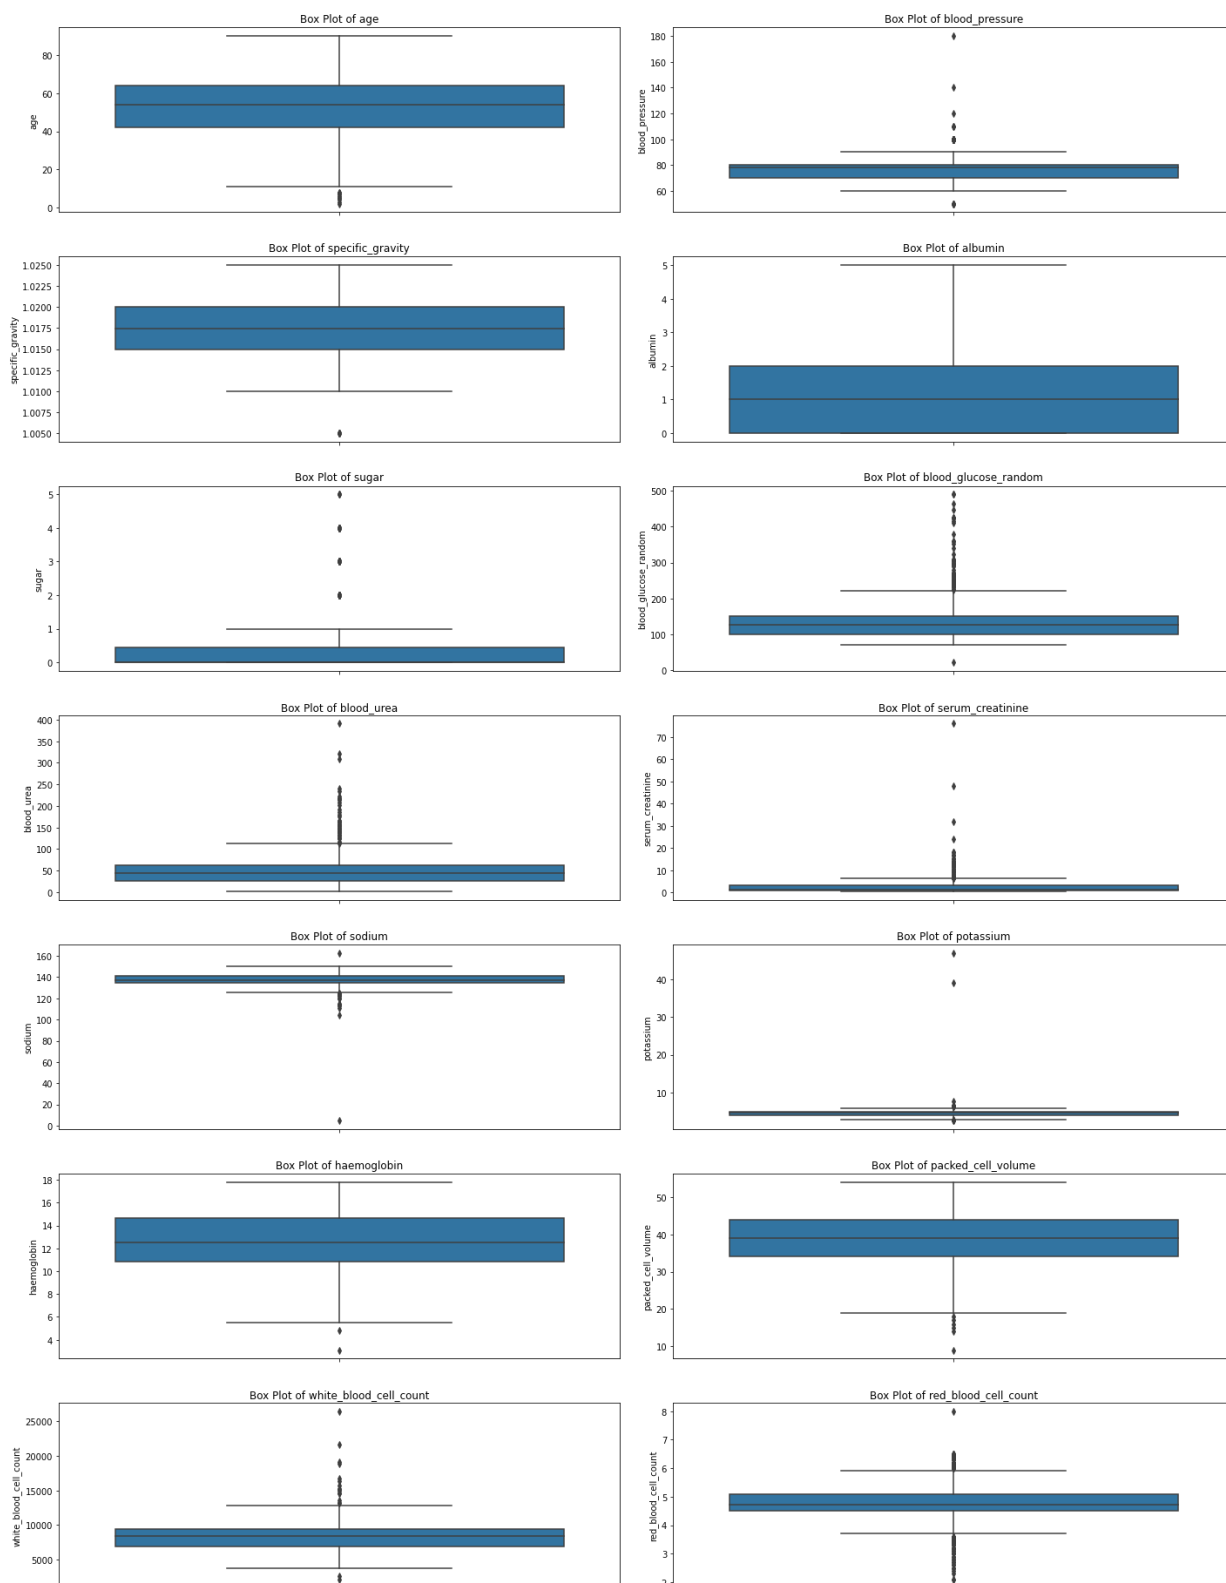
fig, axes = plt.subplots(num_rows, num_cols_per_row, figsize=(20, 26))
fig.tight_layout(pad=3.0) # Adjust the spacing between plots

for i, column in enumerate(num_cols):
    row = i // num_cols_per_row
    col = i % num_cols_per_row


    sns.boxplot(data=df, y=column, ax=axes[row, col])
    axes[row, col].set_title(f'Box Plot of {column}')
    axes[row, col].set_ylabel(column)

# Hide any empty subplots
for i in range(len(num_cols), num_rows * num_cols_per_row):
    row = i // num_cols_per_row
    col = i % num_cols_per_row
    axes[row, col].axis('off')
```

```
plt.show()
```



```
In [41]: correlation_matrix = df[num_cols].corr()
plt.figure(figsize=(16, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Matrix')
plt.show()
```



```
In [44]: target_col = 'class' # target variable

# Calculate the number of rows and columns needed for the grid
num_rows = len(num_cols) // 2 + len(num_cols) % 2
num_cols_per_row = 2

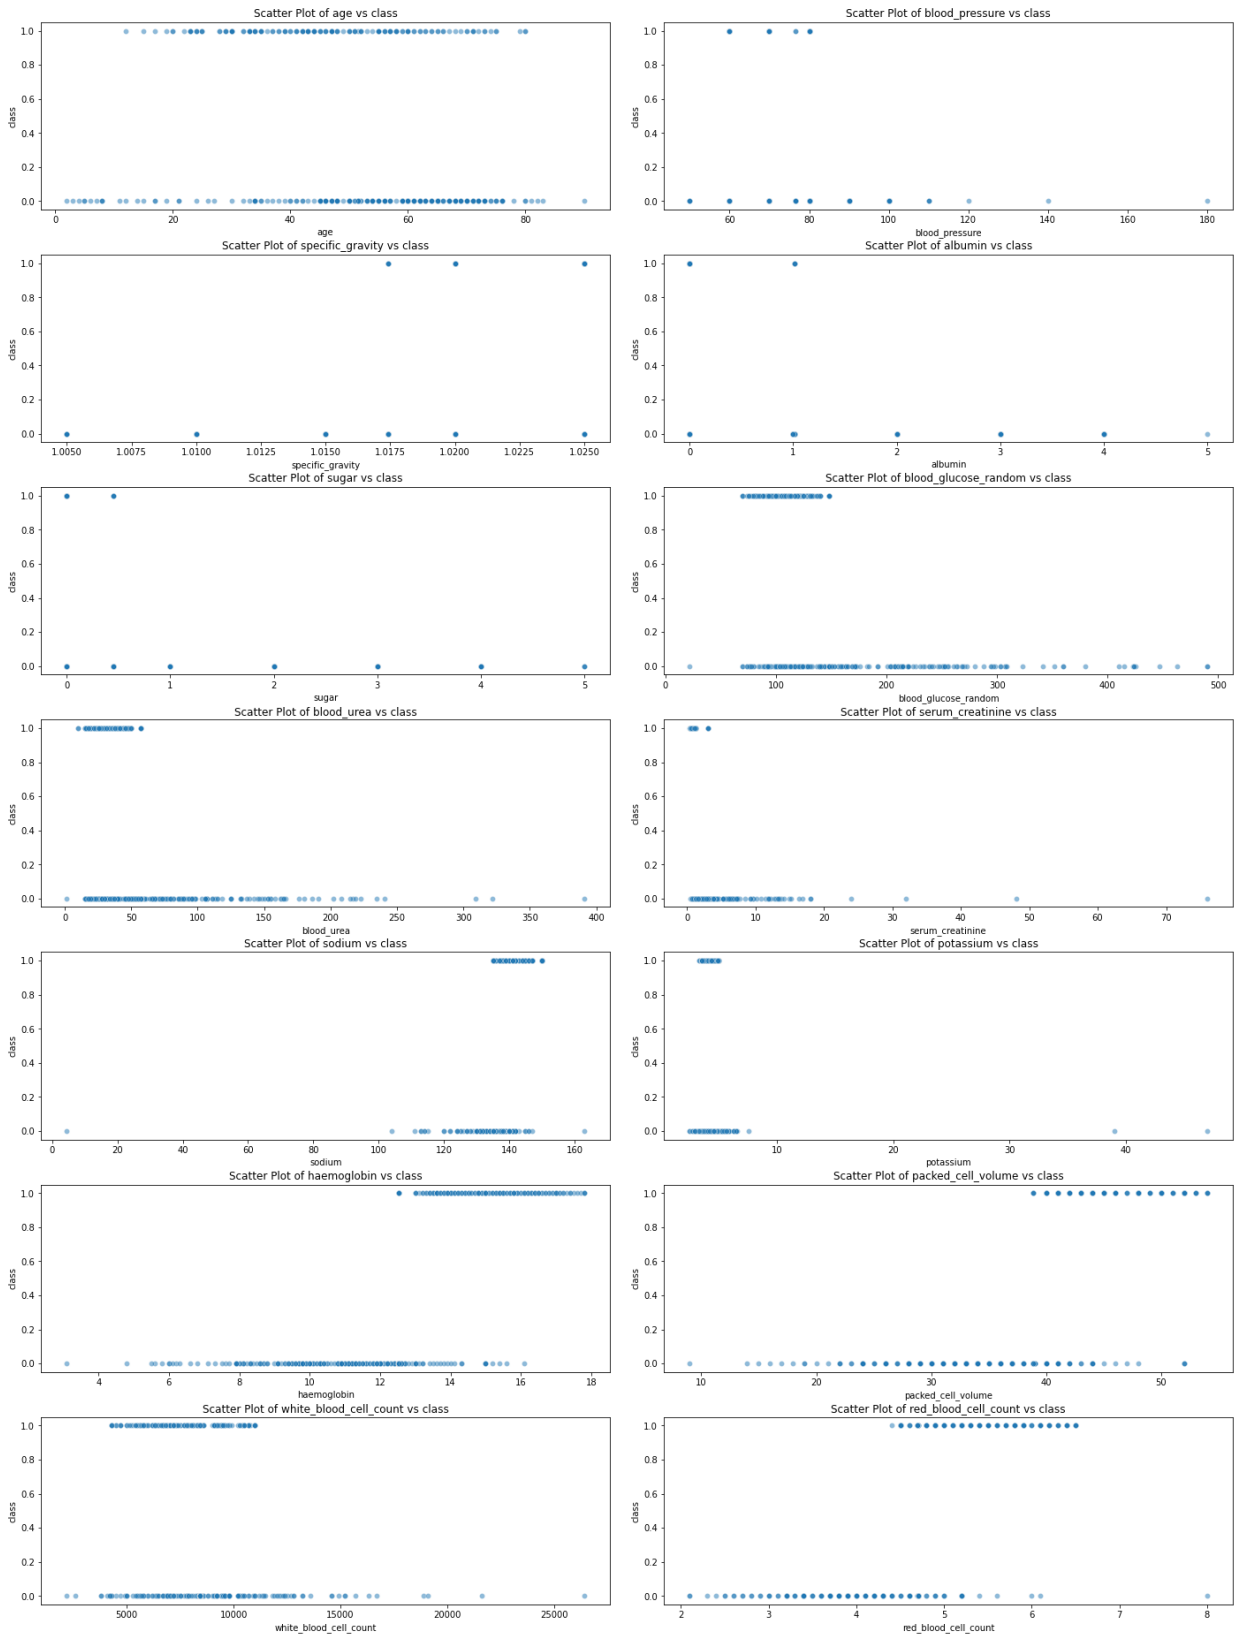
fig, axes = plt.subplots(num_rows, num_cols_per_row, figsize=(20, 26))
fig.tight_layout(pad=3.0) # Adjust the spacing between plots

for i, column in enumerate(num_cols):
    row = i // num_cols_per_row
    col = i % num_cols_per_row

    sns.scatterplot(data=df, x=column, y=target_col, alpha=0.5, ax=axes[row, col])
    axes[row, col].set_title(f'Scatter Plot of {column} vs {target_col}')
    axes[row, col].set_xlabel(column)
    axes[row, col].set_ylabel(target_col)

# Hide any empty subplots
for i in range(len(num_cols), num_rows * num_cols_per_row):
    row = i // num_cols_per_row
    col = i % num_cols_per_row
    axes[row, col].axis('off')

plt.show()
```



```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

Feature Enginnering

```
In [53]: cat_col_unique_counts = {}

for col in cat_cols:
    unique_count = df[col].nunique()
    cat_col_unique_counts[col] = unique_count
    print(f"{col} has: {unique_count} categories\n")

red_blood_cells has: 2 categories

pus_cell has: 2 categories

pus_cell_clumps has: 2 categories

bacteria has: 2 categories

hypertension has: 2 categories

diabetes_mellitus has: 2 categories

coronary_artery_disease has: 2 categories

appetite has: 2 categories

peda_edema has: 2 categories

aanemia has: 2 categories

class has: 2 categories
```

```
In [54]: # LabelEncoding
from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()

for col in cat_cols:
    df[col] = label_encoder.fit_transform(df[col])
```

```
In [55]: df.head(10)
```

Out[55]:

	age	blood_pressure	specific_gravity	albumin	sugar	red_blood_cells	pus_cell	pus_cell_clumps	b
0	48.0	80.000000	1.020	1.0	0.0	1	1	0	
1	7.0	50.000000	1.020	4.0	0.0	1	1	0	
2	62.0	80.000000	1.010	2.0	3.0	1	1	0	
3	48.0	70.000000	1.005	4.0	0.0	1	0	1	
4	51.0	80.000000	1.010	2.0	0.0	1	1	0	
5	60.0	90.000000	1.015	3.0	0.0	1	1	0	
6	68.0	70.000000	1.010	0.0	0.0	1	1	0	
7	24.0	76.469072	1.015	2.0	4.0	1	0	0	
8	52.0	100.000000	1.015	3.0	0.0	1	0	1	
9	53.0	90.000000	1.020	2.0	0.0	0	0	1	

In [56]: `df.isnull().sum()`


Out[56]:

age	0
blood_pressure	0
specific_gravity	0
albumin	0
sugar	0
red_blood_cells	0
pus_cell	0
pus_cell_clumps	0
bacteria	0
blood_glucose_random	0
blood_urea	0
serum_creatinine	0
sodium	0
potassium	0
haemoglobin	0
packed_cell_volume	0
white_blood_cell_count	0
red_blood_cell_count	0
hypertension	0
diabetes_mellitus	0
coronary_artery_disease	0
appetite	0
peda_edema	0
aanemia	0
class	0

dtype: int64

In [57]:

```
plt.figure(figsize=(10, 6))
sns.heatmap(df.isnull(), cmap='viridis', cbar=False)
plt.title('Missing Value Heatmap')
plt.show()
```



In []:

MODEL BUILDING

In []:

Train Test split and Normalization

```
In [58]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.model_selection import train_test_split

# Splitting the data into features (X) and target variable (y)
X = df.drop('class', axis=1) # Features
y = df['class'] # Target variable

# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=

# Create a pipeline with StandardScaler
pipeline_standard = Pipeline([
    ('scaler', StandardScaler()) # Step 1: StandardScaler for normalization
])

# Create a pipeline with MinMaxScaler
```

```

pipeline_minmax = Pipeline([
    ('scaler', MinMaxScaler()) # Step 1: MinMaxScaler for normalization
])

# Fit and transform using the pipeline with StandardScaler
X_train_standard = pipeline_standard.fit_transform(X_train)

# Fit and transform using the pipeline with MinMaxScaler
X_train_minmax = pipeline_minmax.fit_transform(X_train)

```

```

In [73]: print("X shape:", X_train.shape)

print("y shape:", y_train.shape)

```

```

X shape: (320, 24)
y shape: (320,)

```

Model building

Logistic Regression

```

In [82]: from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

param_grid = {
    'solver': ['newton-cg', 'lbfgs', 'liblinear'],
    'penalty': ['l2'], # 'none', 'l1', 'l2', 'elasticnet'
    'C': np.logspace(-5, 2, num=10)
}

grid = GridSearchCV(LogisticRegression(max_iter=1000), param_grid, refit=True, verbose=0)
grid.fit(X_train, y_train)

# Access the best model and its parameters
best_model = grid.best_estimator_
best_params = grid.best_params_
print(f"Best params: {best_params}")

# Evaluate the best model
y_train_pred = best_model.predict(X_train)
y_test_pred = best_model.predict(X_test)

train_accuracy = accuracy_score(y_train, y_train_pred)
test_accuracy = accuracy_score(y_test, y_test_pred)

print(f"Train Accuracy: {train_accuracy:.4f}")
print(f"Test Accuracy: {test_accuracy:.4f}")

# Visualize confusion matrix
conf_matrix = confusion_matrix(y_test, y_test_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')

```



```
plt.ylabel('Actual')
plt.show()


# Print classification report
print(f"Classification Report :- \n {classification_report(y_test, y_test_pred)}")
```

Fitting 5 folds for each of 30 candidates, totalling 150 fits

Best params: {'C': 16.68100537200059, 'penalty': 'l2', 'solver': 'liblinear'}

Train Accuracy: 0.9906

Test Accuracy: 0.9750



Classification Report :-

	precision	recall	f1-score	support
0	0.98	0.98	0.98	52
1	0.96	0.96	0.96	28
accuracy			0.97	80
macro avg	0.97	0.97	0.97	80
weighted avg	0.97	0.97	0.97	80

KNN

```
In [89]: from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

param_grid = {
    'n_neighbors': [3, 5, 7, 9, 11],
    'weights': ['uniform', 'distance'],
    'metric': ['euclidean', 'manhattan']
}

grid = GridSearchCV(KNeighborsClassifier(), param_grid, refit=True, verbose=1, cv=5)
model = grid.fit(X_train, y_train).best_estimator_
```

```
best_params = grid.best_params_  
print(f"Best params: {best_params}")  
  
y_train_pred = model.predict(X_train)  
y_test_pred = model.predict(X_test)  
  
# Define a function to print scores  
def print_score(model, X_train, y_train, X_test, y_test):  
    print("Training Result:\n")  
    acc_train = accuracy_score(y_train, model.predict(X_train))  
    print(f"Training Accuracy Score: {acc_train:.4f}\n")  
    conf_matrix_train = confusion_matrix(y_train, model.predict(X_train))  
  
    print("Testing Result:\n")  
    acc_test = accuracy_score(y_test, model.predict(X_test))  
    print(f"Testing Accuracy Score: {acc_test:.4f}\n")  
    conf_matrix_test = confusion_matrix(y_test, model.predict(X_test))  
  
    print(f"Confusion Matrix (Training):\n{conf_matrix_train}\n")  
    print(f"Confusion Matrix (Testing):\n{conf_matrix_test}\n")  
    print(f"Classification Report (Testing):\n{classification_report(y_test, model.pre  
  
print_score(model, X_train, y_train, X_test, y_test)  
  
# Visualize the confusion matrix  
plt.figure(figsize=(8, 6))  
sns.heatmap(confusion_matrix(y_test, y_test_pred), annot=True, fmt='d', cmap='Blues',  
plt.xlabel('Predicted')  
plt.ylabel('Actual')  
plt.title('Confusion Matrix')  
plt.show()
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits
 Best params: {'metric': 'manhattan', 'n_neighbors': 3, 'weights': 'distance'}
 Training Result:

Training Accuracy Score: 1.0000

Testing Result:

Testing Accuracy Score: 0.7875

Confusion Matrix (Training):


```
[[198  0]
 [  0 122]]
```

Confusion Matrix (Testing):

```
[[36 16]
 [ 1 27]]
```

Classification Report (Testing):

	precision	recall	f1-score	support
0	0.97	0.69	0.81	52
1	0.63	0.96	0.76	28
accuracy			0.79	80
macro avg	0.80	0.83	0.78	80
weighted avg	0.85	0.79	0.79	80



Random Forest

```
In [97]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

rd_clf = RandomForestClassifier(criterion='entropy', max_depth=11, max_features='auto')
```

```

rd_clf.fit(X_train, y_train)

# Accuracy score, confusion matrix and classification report of random forest
rd_clf_acc = accuracy_score(y_test, rd_clf.predict(X_test))

print(f"Training Accuracy of Random Forest Classifier is {accuracy_score(y_train, rd_clf.predict(X_train))}")
print(f"Test Accuracy of Random Forest Classifier is {rd_clf_acc} \n")

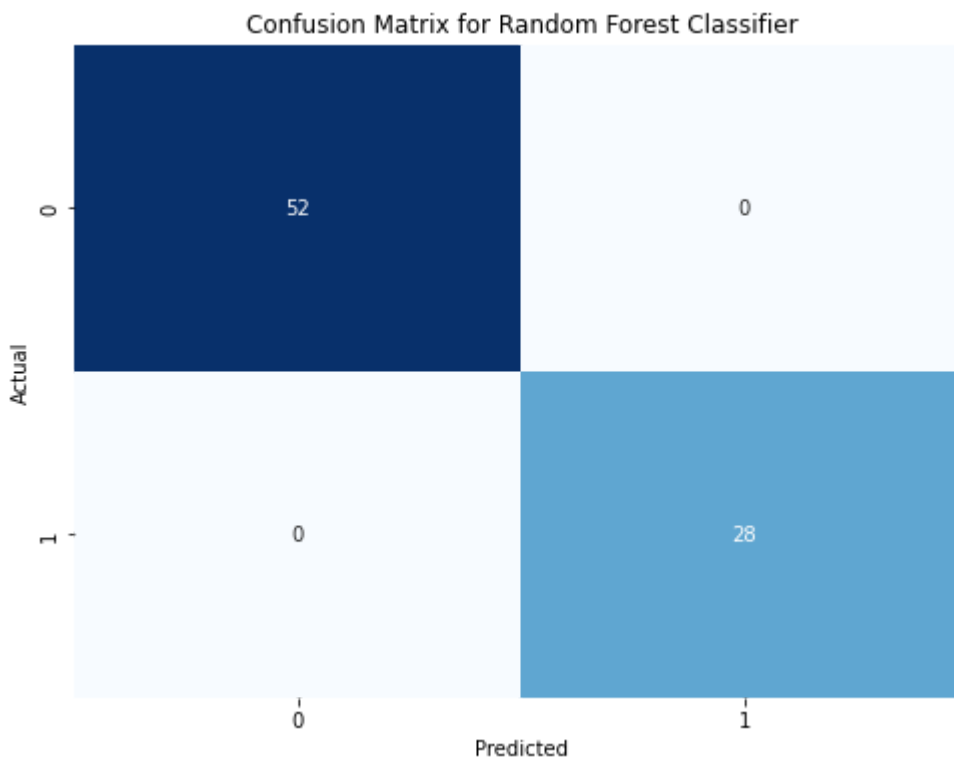
#conf_matrix_rf = confusion_matrix(y_test, rd_clf.predict(X_test))
# Visualize the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_rf, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for Random Forest Classifier')
plt.show()

#print(f"Confusion Matrix :- \n{conf_matrix_rf}\n")
print(f"Classification Report :- \n {classification_report(y_test, rd_clf.predict(X_test))}")

```

Training Accuracy of Random Forest Classifier is 1.0

Test Accuracy of Random Forest Classifier is 0.9875



Classification Report :-

	precision	recall	f1-score	support
0	1.00	0.98	0.99	52
1	0.97	1.00	0.98	28
accuracy			0.99	80
macro avg	0.98	0.99	0.99	80
weighted avg	0.99	0.99	0.99	80

SVM

```
In [96]: from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Create an instance of the SVM classifier
svm_clf = SVC(kernel='rbf', C=1, gamma='auto')

# Train the SVM model
svm_clf.fit(X_train, y_train)

# Make predictions
y_train_pred = svm_clf.predict(X_train)
y_test_pred = svm_clf.predict(X_test)

# Calculate accuracy scores
svm_train_acc = accuracy_score(y_train, y_train_pred)
svm_test_acc = accuracy_score(y_test, y_test_pred)


print(f"Training Accuracy of SVM Classifier: {svm_train_acc:.4f}")
print(f"Test Accuracy of SVM Classifier: {svm_test_acc:.4f}")

# Calculate confusion matrix
#conf_matrix_svm = confusion_matrix(y_test, y_test_pred)

# Visualize the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_svm, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for SVM Classifier')
plt.show()

#print(f"Confusion Matrix:\n{conf_matrix_svm}")
print(f"Classification Report:\n{classification_report(y_test, y_test_pred)}")
```

Training Accuracy of SVM Classifier: 1.0000
Test Accuracy of SVM Classifier: 0.6500



Classification Report:

	precision	recall	f1-score	support
0	0.65	1.00	0.79	52
1	0.00	0.00	0.00	28
accuracy			0.65	80
macro avg	0.33	0.50	0.39	80
weighted avg	0.42	0.65	0.51	80

Naive Biyes

```
In [100... from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns

# Create an instance of the Naive Bayes classifier
nb_clf = GaussianNB()

# Train the Naive Bayes model
nb_clf.fit(X_train, y_train)

# Make predictions
y_train_pred = nb_clf.predict(X_train)
y_test_pred = nb_clf.predict(X_test)

# Calculate accuracy scores
nb_train_acc = accuracy_score(y_train, y_train_pred)
nb_test_acc = accuracy_score(y_test, y_test_pred)

print(f"Training Accuracy of Naive Bayes Classifier: {nb_train_acc:.4f}")
print(f"Test Accuracy of Naive Bayes Classifier: {nb_test_acc:.4f}")
```

```
# Calculate confusion matrix
conf_matrix_nb = confusion_matrix(y_test, y_test_pred)

print(f"Confusion Matrix:\n{conf_matrix_nb}")

# Visualize the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_nb, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for Naive Bayes Classifier')
plt.show()

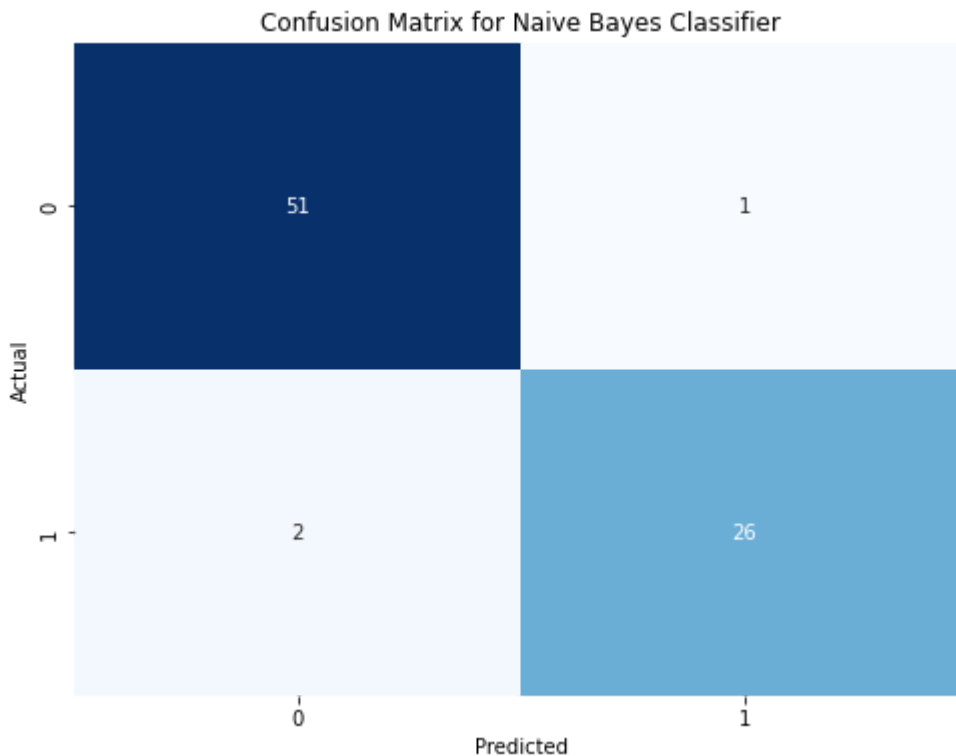
print(f"Classification Report:\n{classification_report(y_test, y_test_pred)}")
```

Training Accuracy of Naive Bayes Classifier: 0.9625

Test Accuracy of Naive Bayes Classifier: 0.9625

Confusion Matrix:

```
[[51  1]
 [ 2 26]]
```



Classification Report:

	precision	recall	f1-score	support
0	0.96	0.98	0.97	52
1	0.96	0.93	0.95	28
accuracy			0.96	80
macro avg	0.96	0.95	0.96	80
weighted avg	0.96	0.96	0.96	80

KNN

In [102...

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
# Create an instance of the KNN classifier
knn_clf = KNeighborsClassifier(n_neighbors=5) # You can adjust the number of neighbors

# Train the KNN model
knn_clf.fit(X_train, y_train)

# Make predictions
y_train_pred = knn_clf.predict(X_train)
y_test_pred = knn_clf.predict(X_test)

# Calculate accuracy scores
knn_train_acc = accuracy_score(y_train, y_train_pred)
knn_test_acc = accuracy_score(y_test, y_test_pred)

print(f"Training Accuracy of KNN Classifier: {knn_train_acc:.4f}")
print(f"Test Accuracy of KNN Classifier: {knn_test_acc:.4f}")

# Calculate confusion matrix
conf_matrix_knn = confusion_matrix(y_test, y_test_pred)

print(f"Confusion Matrix:\n{conf_matrix_knn}")
# Visualize the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_knn, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for KNN Classifier')
plt.show()

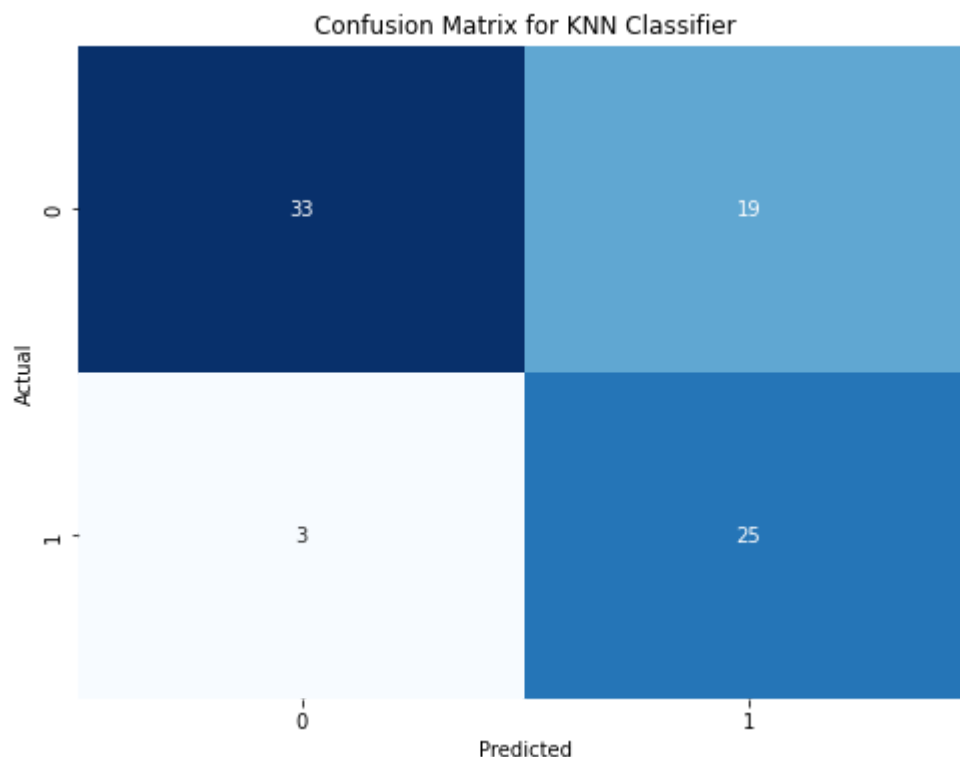
print(f"Classification Report:\n{classification_report(y_test, y_test_pred)}")
```

Training Accuracy of KNN Classifier: 0.8531

Test Accuracy of KNN Classifier: 0.7250

Confusion Matrix:

```
[[33 19]
 [ 3 25]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.63	0.75	52
1	0.57	0.89	0.69	28
accuracy			0.73	80
macro avg	0.74	0.76	0.72	80
weighted avg	0.79	0.72	0.73	80

In [107...

```
from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_neighbors': [3, 5, 7, 9, 11],
    'weights': ['uniform', 'distance'],
    'metric': ['euclidean', 'manhattan']
}

grid = GridSearchCV(KNeighborsClassifier(), param_grid, refit=True, verbose=1, cv=5)
best_knn_model = grid.fit(X_train, y_train).best_estimator_
from sklearn.metrics import accuracy_score

# Predicting on the test set
y_test_pred = best_knn_model.predict(X_test)

# Calculating accuracy
knn_accuracy = accuracy_score(y_test, y_test_pred)
print(f"Test Accuracy of the Best KNN Model: {knn_accuracy:.4f}")
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits
 Test Accuracy of the Best KNN Model: 0.7875

Logistic Regression

In [104...

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Create an instance of the Logistic Regression classifier
logreg_clf = LogisticRegression()

# Train the Logistic Regression model
logreg_clf.fit(X_train, y_train)

# Make predictions
y_train_pred = logreg_clf.predict(X_train)
y_test_pred = logreg_clf.predict(X_test)

# Calculate accuracy scores
logreg_train_acc = accuracy_score(y_train, y_train_pred)
logreg_test_acc = accuracy_score(y_test, y_test_pred)

print(f"Training Accuracy of Logistic Regression Classifier: {logreg_train_acc:.4f}")
print(f"Test Accuracy of Logistic Regression Classifier: {logreg_test_acc:.4f}")

# Calculate confusion matrix
conf_matrix_logreg = confusion_matrix(y_test, y_test_pred)

print(f"Confusion Matrix:\n{conf_matrix_logreg}")
# Visualize the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_logreg, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for Logistic Regression Classifier')
plt.show()

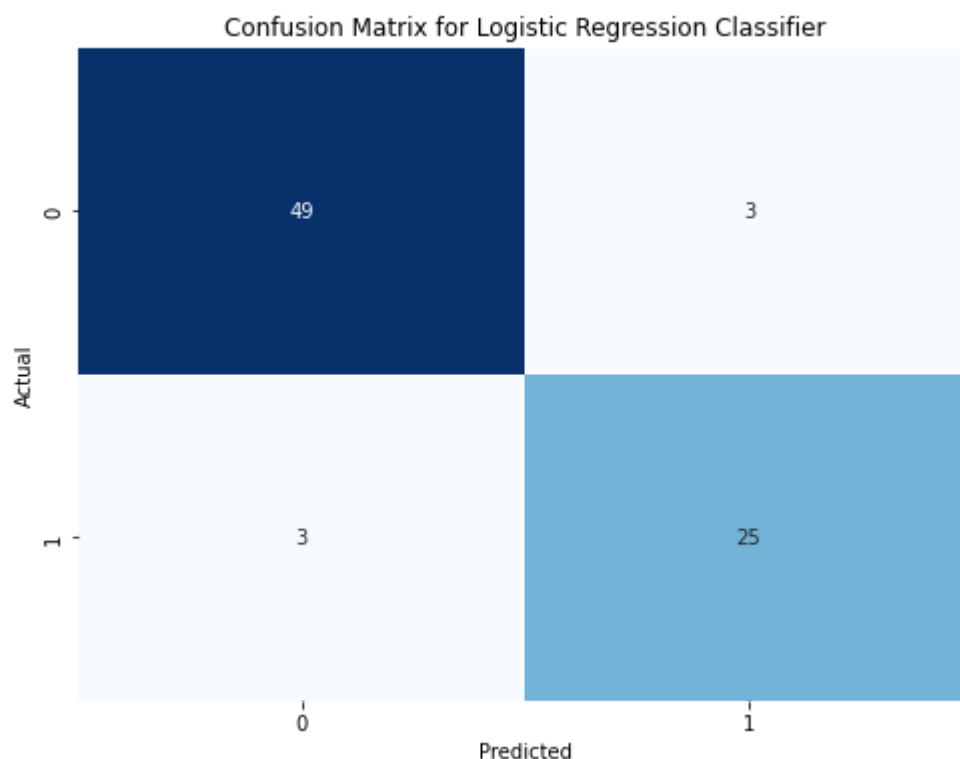
print(f"Classification Report:\n{classification_report(y_test, y_test_pred)}")
```

Training Accuracy of Logistic Regression Classifier: 0.9187

Test Accuracy of Logistic Regression Classifier: 0.9250

Confusion Matrix:

```
[[49  3]
 [ 3 25]]
```



Classification Report:

	precision	recall	f1-score	support
0	0.94	0.94	0.94	52
1	0.89	0.89	0.89	28
accuracy			0.93	80
macro avg	0.92	0.92	0.92	80
weighted avg	0.93	0.93	0.93	80

In [105...

```

from sklearn.model_selection import GridSearchCV

param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100],
    'penalty': ['l1', 'l2'],
    'solver': ['liblinear', 'saga']
}

grid_search = GridSearchCV(LogisticRegression(), param_grid, cv=5)
grid_search.fit(X_train, y_train)

best_params = grid_search.best_params_
best_model = grid_search.best_estimator_

print("Best Hyperparameters:", best_params)

# Evaluate the best model
y_test_pred = best_model.predict(X_test)
logreg_test_acc = accuracy_score(y_test, y_test_pred)

print(f"Test Accuracy of Best Logistic Regression Classifier: {logreg_test_acc:.4f}")

Best Hyperparameters: {'C': 100, 'penalty': 'l2', 'solver': 'liblinear'}
Test Accuracy of Best Logistic Regression Classifier: 0.9750

```

Forward Neural Network

```
In [109... import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from keras.models import Sequential
from keras.layers import Dense

In [110... # Splitting the data into features (X) and target variable (y)
X = df.drop('class', axis=1) # Features
y = df['class']             # Target variable

# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=

# Standardize the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

In [111... model = Sequential()
model.add(Dense(64, activation='relu', input_dim=X_train_scaled.shape[1]))
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

In [112... model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

In [113... history = model.fit(X_train_scaled, y_train, epochs=10, batch_size=32, validation_split=
```

```

Epoch 1/10
8/8 [=====] - 1s 39ms/step - loss: 0.5139 - accuracy: 0.7109
- val_loss: 0.4125 - val_accuracy: 0.9062
Epoch 2/10
8/8 [=====] - 0s 6ms/step - loss: 0.3431 - accuracy: 0.9766
- val_loss: 0.2877 - val_accuracy: 0.9844
Epoch 3/10
8/8 [=====] - 0s 5ms/step - loss: 0.2475 - accuracy: 0.9844
- val_loss: 0.2118 - val_accuracy: 0.9844
Epoch 4/10
8/8 [=====] - 0s 5ms/step - loss: 0.1837 - accuracy: 0.9844
- val_loss: 0.1592 - val_accuracy: 0.9688
Epoch 5/10
8/8 [=====] - 0s 6ms/step - loss: 0.1360 - accuracy: 0.9883
- val_loss: 0.1230 - val_accuracy: 0.9844
Epoch 6/10
8/8 [=====] - 0s 5ms/step - loss: 0.1014 - accuracy: 0.9922
- val_loss: 0.0979 - val_accuracy: 1.0000
Epoch 7/10
8/8 [=====] - 0s 5ms/step - loss: 0.0778 - accuracy: 0.9961
- val_loss: 0.0800 - val_accuracy: 1.0000
Epoch 8/10
8/8 [=====] - 0s 5ms/step - loss: 0.0608 - accuracy: 0.9961
- val_loss: 0.0672 - val_accuracy: 1.0000
Epoch 9/10
8/8 [=====] - 0s 5ms/step - loss: 0.0486 - accuracy: 0.9961
- val_loss: 0.0578 - val_accuracy: 1.0000
Epoch 10/10
8/8 [=====] - 0s 5ms/step - loss: 0.0390 - accuracy: 0.9961
- val_loss: 0.0511 - val_accuracy: 1.0000

```

In [114...

```

test_loss, test_accuracy = model.evaluate(X_test_scaled, y_test)
print(f"Test Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_accuracy:.4f}")

```

```

3/3 [=====] - 0s 2ms/step - loss: 0.0431 - accuracy: 0.9875
Test Loss: 0.0431
Test Accuracy: 0.9875

```

In []:

model evaluation

In [120...

```

from sklearn.linear_model import LogisticRegression

# Create and train a Logistic Regression model
lr_model = LogisticRegression()
lr_model.fit(X_train, y_train)

# Evaluate Logistic Regression
lr_pred = lr_model.predict(X_test)
print_metrics("Logistic Regression", y_test, lr_pred)

```

Logistic Regression Metrics:

Accuracy: 0.9250

Precision: 0.8929

Recall: 0.8929

F1-score: 0.8929

ROC AUC: 0.9176

In [123...

```

#from sklearn.linear_model import LogisticRegression
#from sklearn.ensemble import RandomForestClassifier
#from sklearn.svm import SVC
#from sklearn.naive_bayes import GaussianNB
#from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, r

# Create an empty DataFrame to store results
results_df = pd.DataFrame(columns=['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score', 'ROC AUC'])

# List of models
models = [
    ('Logistic Regression', LogisticRegression()),
    ('Random Forest', RandomForestClassifier()),
    ('Support Vector Machine', SVC()),
    ('Naive Bayes', GaussianNB()),
    ('K-Nearest Neighbors', KNeighborsClassifier())
]

# Iterate through each model
for name, model in models:
    model.fit(X_train, y_train) # Train the model
    y_pred = model.predict(X_test) # Predict on the test set

    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    roc_auc = roc_auc_score(y_test, y_pred)

    results_df = results_df.append({
        'Model': name,
        'Accuracy': accuracy,
        'Precision': precision,
        'Recall': recall,
        'F1 Score': f1,
        'ROC AUC': roc_auc
    }, ignore_index=True)

# Display the results DataFrame
print(results_df)

```

	Model	Accuracy	Precision	Recall	F1 Score	ROC AUC
0	Logistic Regression	0.9250	0.8929	0.8929	0.8929	0.9176
1	Random Forest	1.0000	1.0000	1.0000	1.0000	1.0000
2	Support Vector Machine	0.6500	0.0000	0.0000	0.0000	0.5000
3	Naive Bayes	0.9625	0.9630	0.9286	0.9455	0.9547
4	K-Nearest Neighbors	0.7250	0.5682	0.8929	0.6944	0.7637

In [125...

```

#import matplotlib.pyplot as plt

# Define colors for each model

```

```

colors = ['skyblue', 'lightgreen', 'lightcoral', 'orange', 'purple']


# Sort the results DataFrame by Accuracy Score in descending order
results_df_sorted = results_df.sort_values(by='Accuracy', ascending=False)

# Create a bar chart with custom colors
plt.figure(figsize=(10, 6))
plt.barh(results_df_sorted['Model'], results_df_sorted['Accuracy'], color=colors)
plt.xlabel('Accuracy')
plt.ylabel('Model')
plt.title('Top Models by Accuracy')
plt.gca().invert_yaxis() # Invert y-axis to have the highest accuracy on top

# Add a Legend with model names and corresponding colors
plt.legend(results_df_sorted['Model'], loc='lower right')

plt.show()

```



In [127...

```

# eady trained the Random Forest model
rd_clf.fit(X_train, y_train)

# Get feature importances from the trained Random Forest model
feature_importances = rd_clf.feature_importances_

# Create a DataFrame with feature names and their corresponding importances
feature_importance_df = pd.DataFrame({'Feature': X.columns, 'Importance': feature_importances})

# Sort the DataFrame by importance in descending order
feature_importance_df_sorted = feature_importance_df.sort_values(by='Importance', ascending=False)


# Choose different colors for the bars
colors = ['skyblue', 'salmon', 'lightgreen', 'orange', 'lightcoral', 'lightskyblue', 'lightpink']

# Plot the top N features and their importances with different colors
top_n = 10 # Number of top features to display
plt.figure(figsize=(10, 6))
plt.barh(feature_importance_df_sorted['Feature'][:top_n], feature_importance_df_sorted['Importance'][:top_n], color=colors[:top_n])
plt.xlabel('Importance')

```

```
plt.ylabel('Feature')
plt.title(f'Top {top_n} Features by Importance')
plt.gca().invert_yaxis() # Invert y-axis to have the highest importance on top

plt.show()
```



Model Save

In [130...

```
import joblib

# List of models
models = [
    ('Logistic Regression', LogisticRegression()),
    ('Random Forest', RandomForestClassifier()),
    ('Support Vector Machine', SVC()),
    ('Naive Bayes', GaussianNB()),
    ('K-Nearest Neighbors', KNeighborsClassifier())
]

# Save each model
for name, model in models:
    model_filename = f'{name.lower().replace(" ", "_")}_model.joblib'
    joblib.dump(model, model_filename)
    print(f"Model '{name}' saved as '{model_filename}'")
```

Model 'Logistic Regression' saved as 'logistic_regression_model.joblib'

Model 'Random Forest' saved as 'random_forest_model.joblib'

Model 'Support Vector Machine' saved as 'support_vector_machine_model.joblib'

Model 'Naive Bayes' saved as 'naive_bayes_model.joblib'

Model 'K-Nearest Neighbors' saved as 'k-nearest_neighbors_model.joblib'

In [131...

```
import pickle

# List of models
models = [
    ('Logistic Regression', LogisticRegression()),
    ('Random Forest', RandomForestClassifier()),
    ('Support Vector Machine', SVC()),
```



```
('Naive Bayes', GaussianNB()),  
('K-Nearest Neighbors', KNeighborsClassifier())  
]  
# Save each model  
for name, model in models:  
    model_filename = f'{name.lower().replace(" ", "_")}_model.pkl'  
    with open(model_filename, 'wb') as file:  
        pickle.dump(model, file)  
    print(f"Model '{name}' saved as '{model_filename}'")
```

Model 'Logistic Regression' saved as 'logistic_regression_model.pkl'
Model 'Random Forest' saved as 'random_forest_model.pkl'
Model 'Support Vector Machine' saved as 'support_vector_machine_model.pkl'
Model 'Naive Bayes' saved as 'naive_bayes_model.pkl'
Model 'K-Nearest Neighbors' saved as 'k-nearest_neighbors_model.pkl'

In [132...

```
import pickle  
  
# Create a dictionary to store the models  
all_models = {}  
for name, model in models:  
    all_models[name] = model  
  
# Save the dictionary of models  
with open('all_models.pkl', 'wb') as file:  
    pickle.dump(all_models, file)  
  
print("All models saved as 'all_models.pkl'")
```

All models saved as 'all_models.pkl'

In []: